An Extensible Software Transport Layer for GPU Networking

Yang Zhou

with: Zhongjie Chen, Ziming Mao, ChonLam Lao, Shuo Yang, Pravein Govindan Kannan, Jiaqi Gao, Yilong Zhao, Yongji Wu, Kaichao You, Fengyuan Ren, Zhiying Xu, Costin Raiciu, Ion Stoica

tinyurl.com/uccl-paper ⋄ github.com/uccl-project/uccl

June 2025

About UCCL Project

Building the fastest collective communication library (CCL)

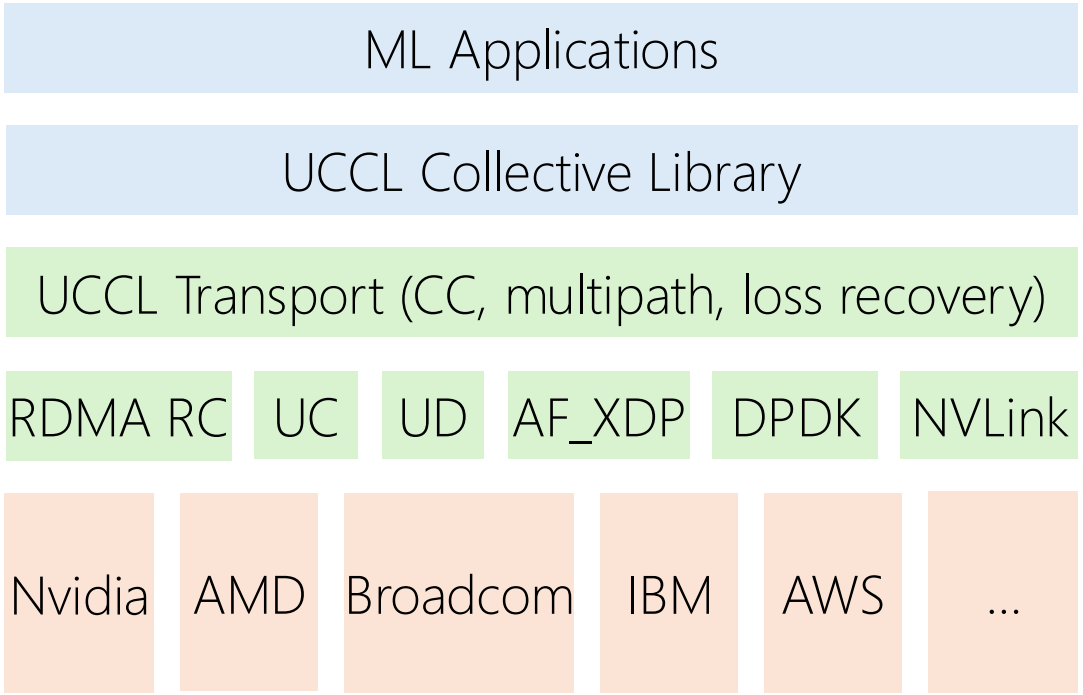

Rearchitecting CCL layer (keep NCCL API)

Rearchitecting transport layer (this talk)

- Up to 3.3× higher tput over NCCL on AWS

Supporting heterogeneous vendors

Open and collaborative platform: github.com/uccl-project/uccl

Fast Evolving ML Workloads

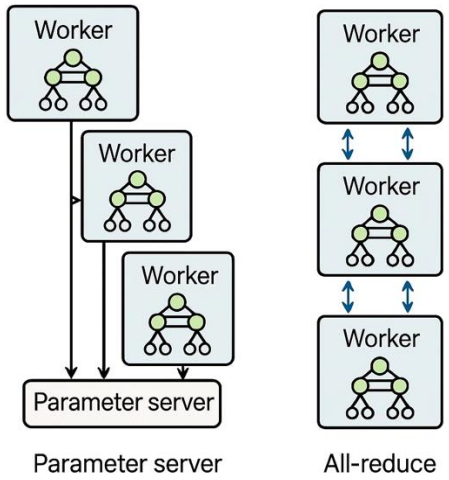

DNN training: parameter servers, allreduce

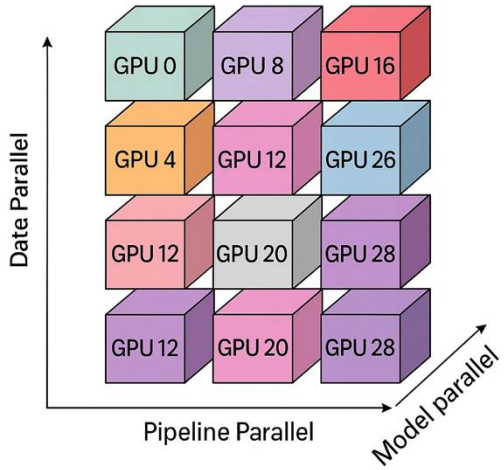

LLM training: allreduce, allgather, and reduce-scatter

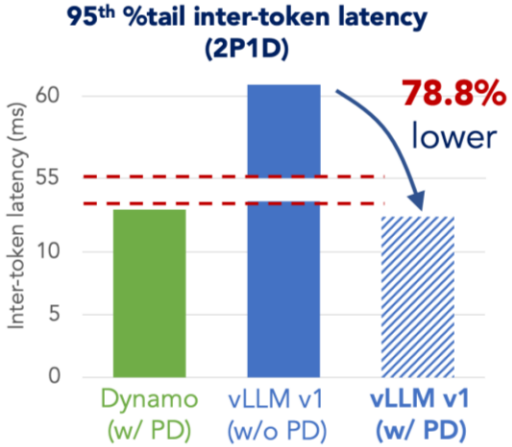

PD disaggregation: P2P transfer

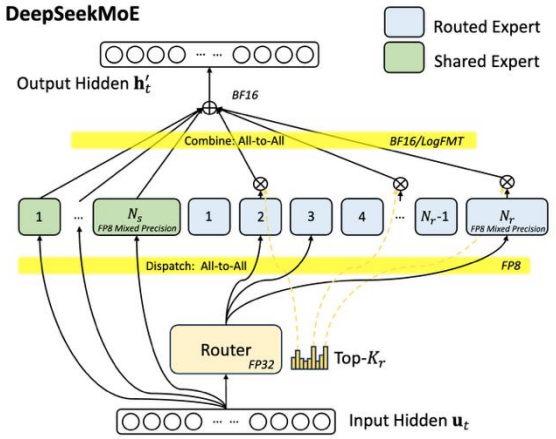

DeepSeek-V3 MoE: all-to-all like

~2015 ~2020 2024 2025 Time

Slowly Evolving Networking

Host transport on RDMA NICs is hard to adapt to better suit ML workloads

- Hardware changes are time-consuming

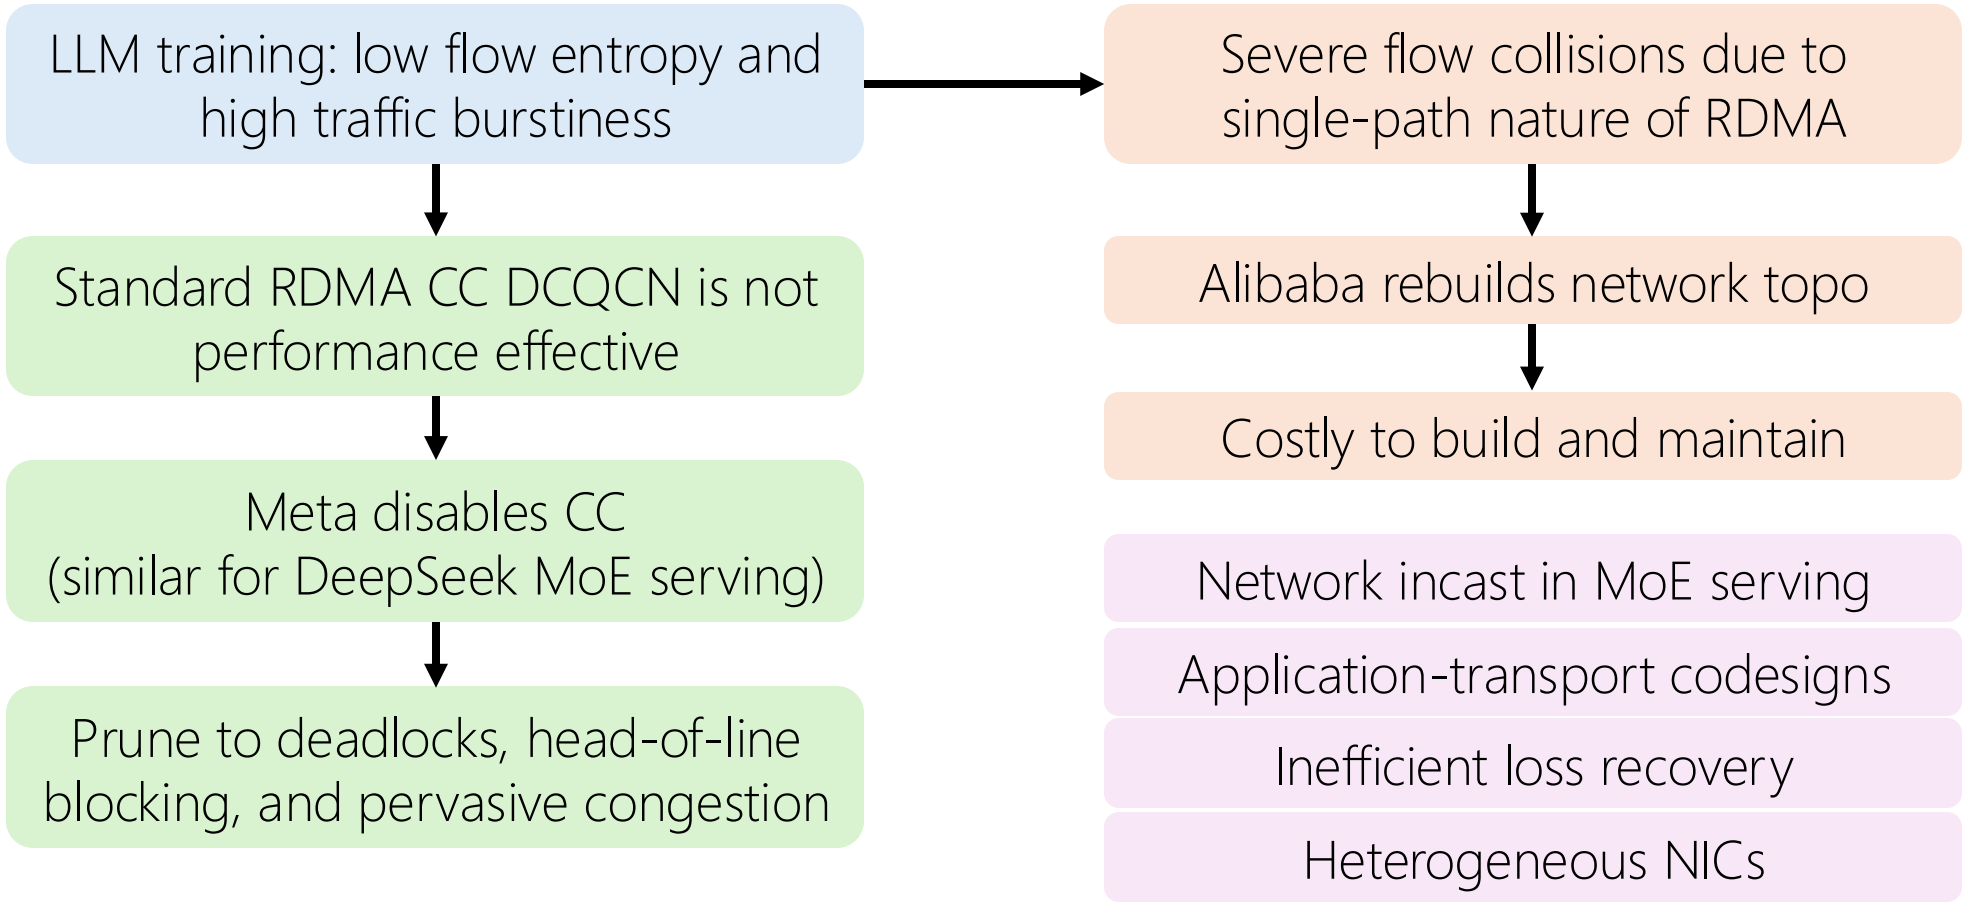

Overarching Problem: Network Extensibility

UCCL approach: a software-only extensible transport for GPU networking

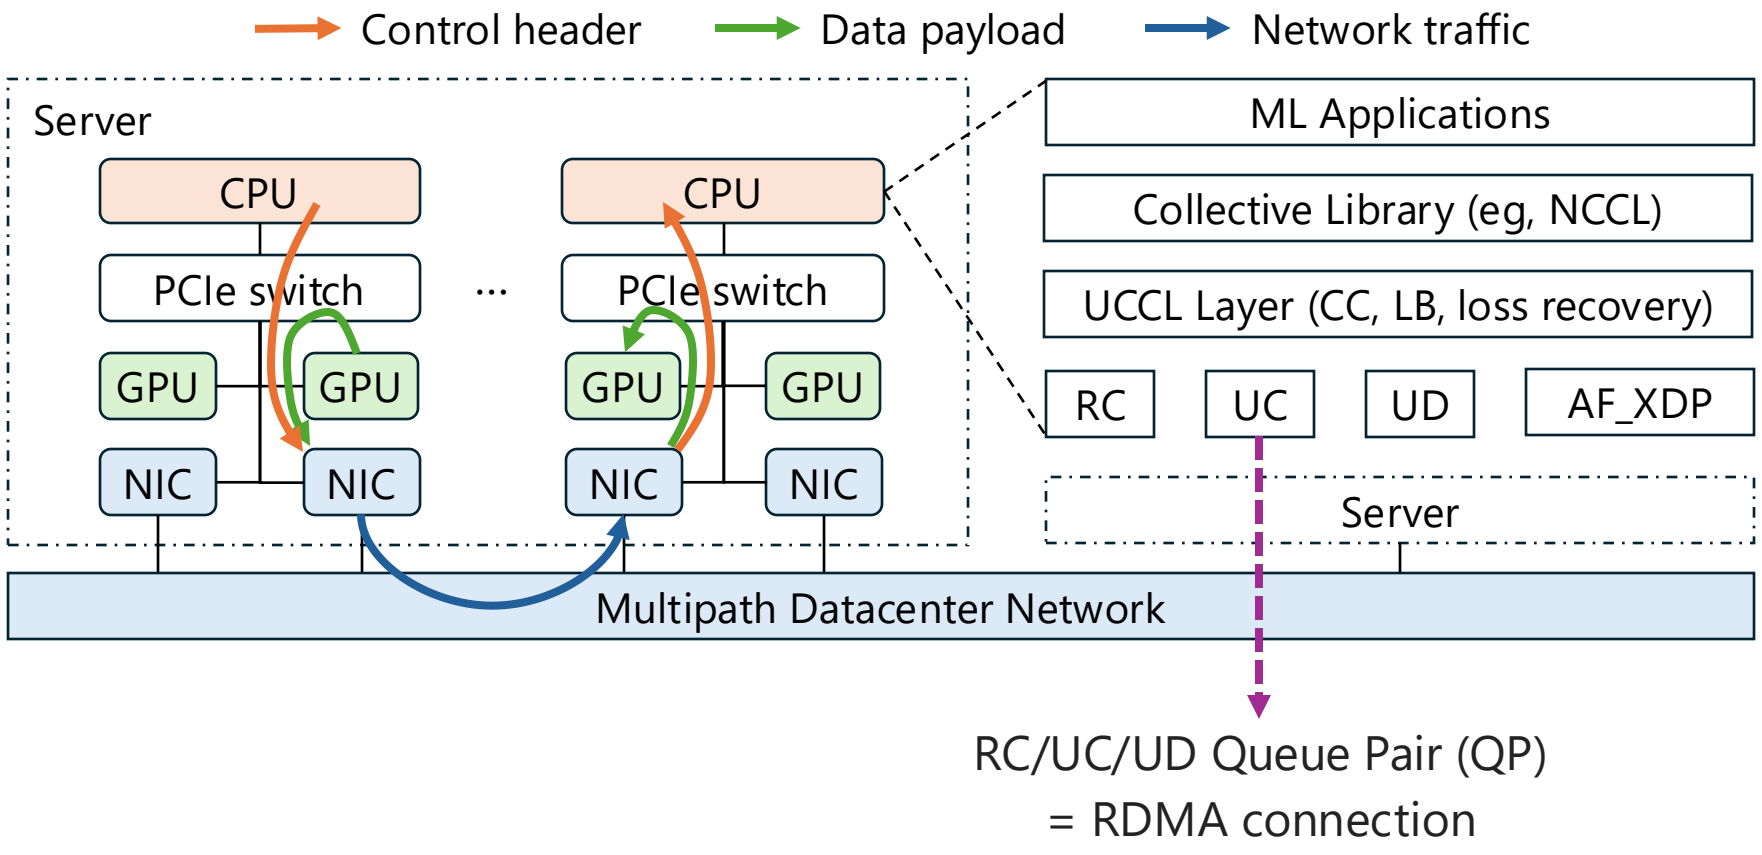

UCCL Key Challenges

- How to **decouple** the data and control paths for existing RDMA NICs?
  - Eg, Nvidia NICs, Broadcom NICs, AWS EFA NICs

- How to achieve hardware-level **performance** for software control path?
  - Eg, 3.2 Tbps inter-server bandwidth

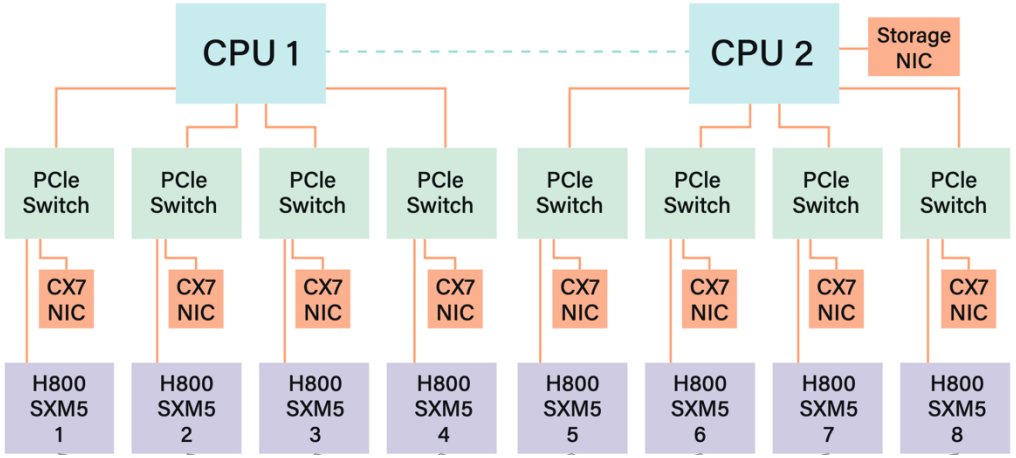

Technique #1: Decoupling Control & Data Path

Leveraging UC/RC QPs + RDMA write with immediate

- Eg, for Nvidia and Broadcom NICs (that support UC or allow disabling RC's CC logic)

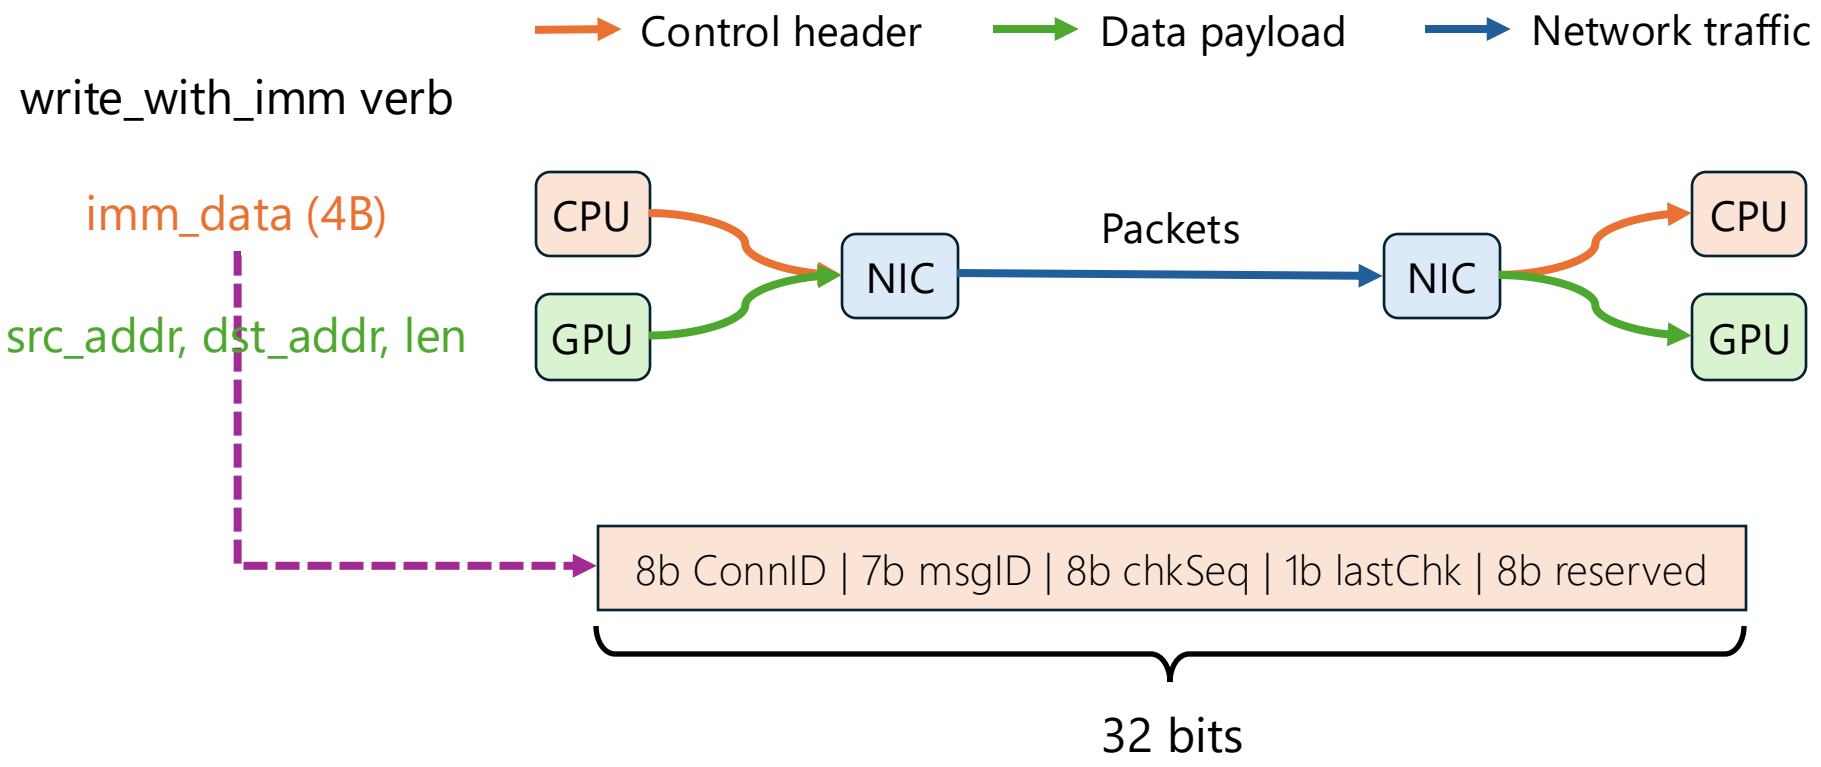

Technique #1: Decoupling Control & Data Path

Leveraging UD QPs + send/recv with scatter-gather list

- Eg, for AWS EFA NICs (that cannot disable RC's CC logic)

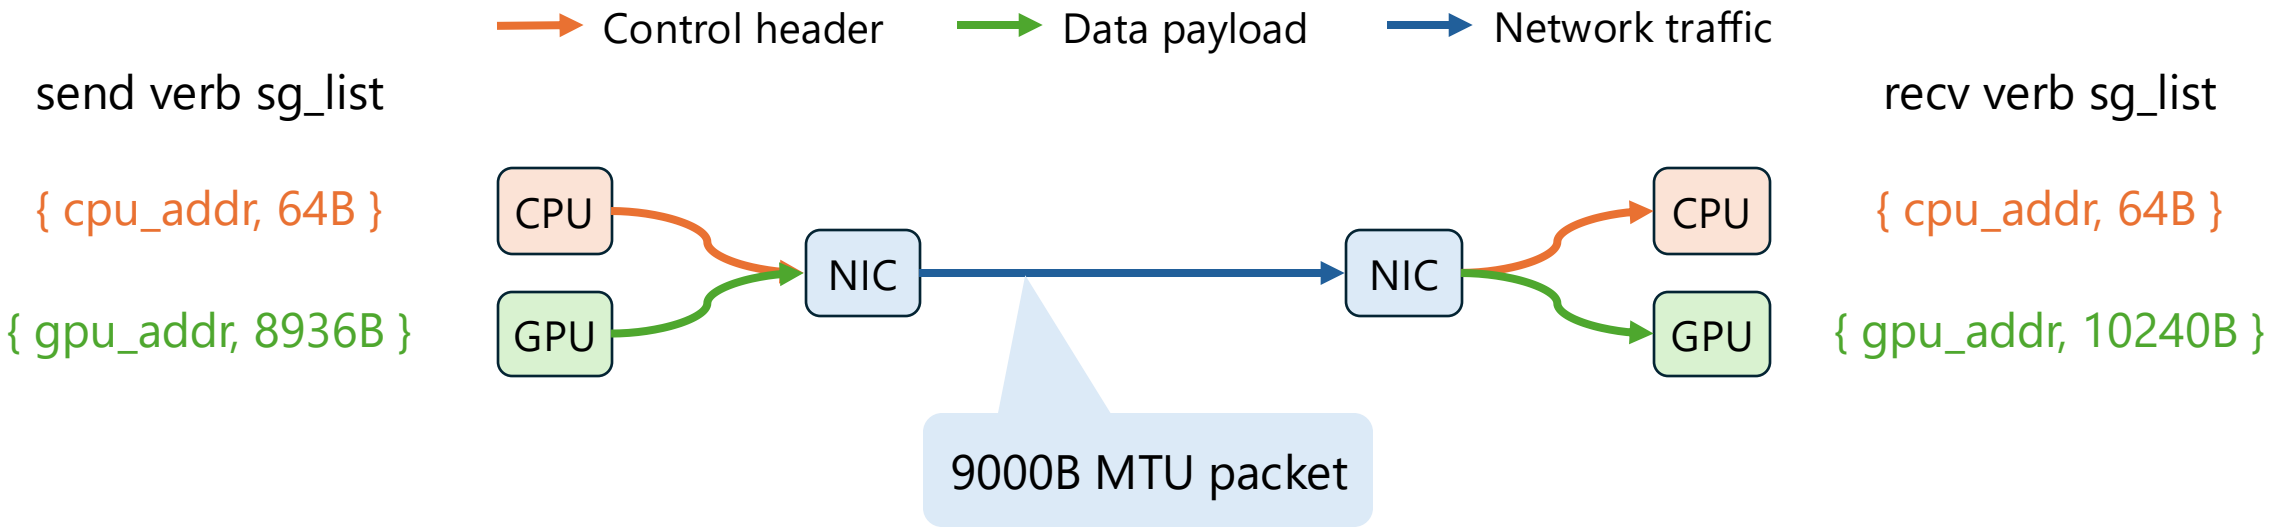

Technique #1: Decoupling Control & Data Path

Leveraging UD QPs + send/recv with scatter-gather list

- Eg, for AWS EFA NICs (that cannot disable RC's CC logic)

Handling out-of-order packet delivery

- Fusing scattered memcpy at the receiver GPU

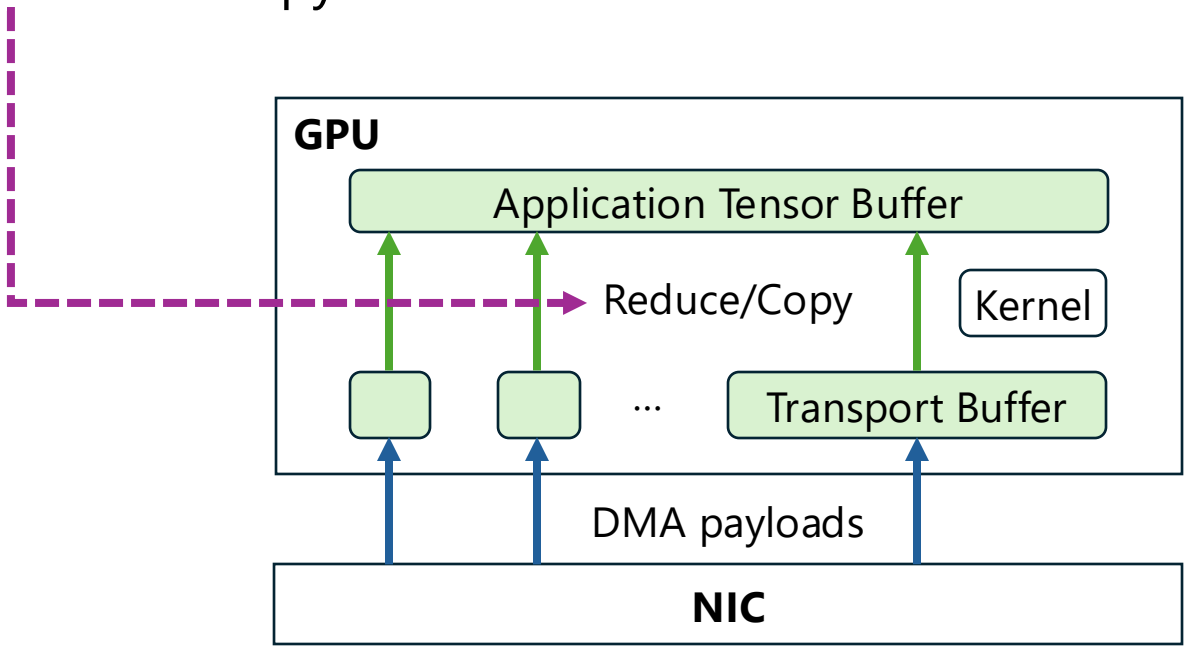

Technique #1: Decoupling Control & Data Path

Multipathing with packet spraying

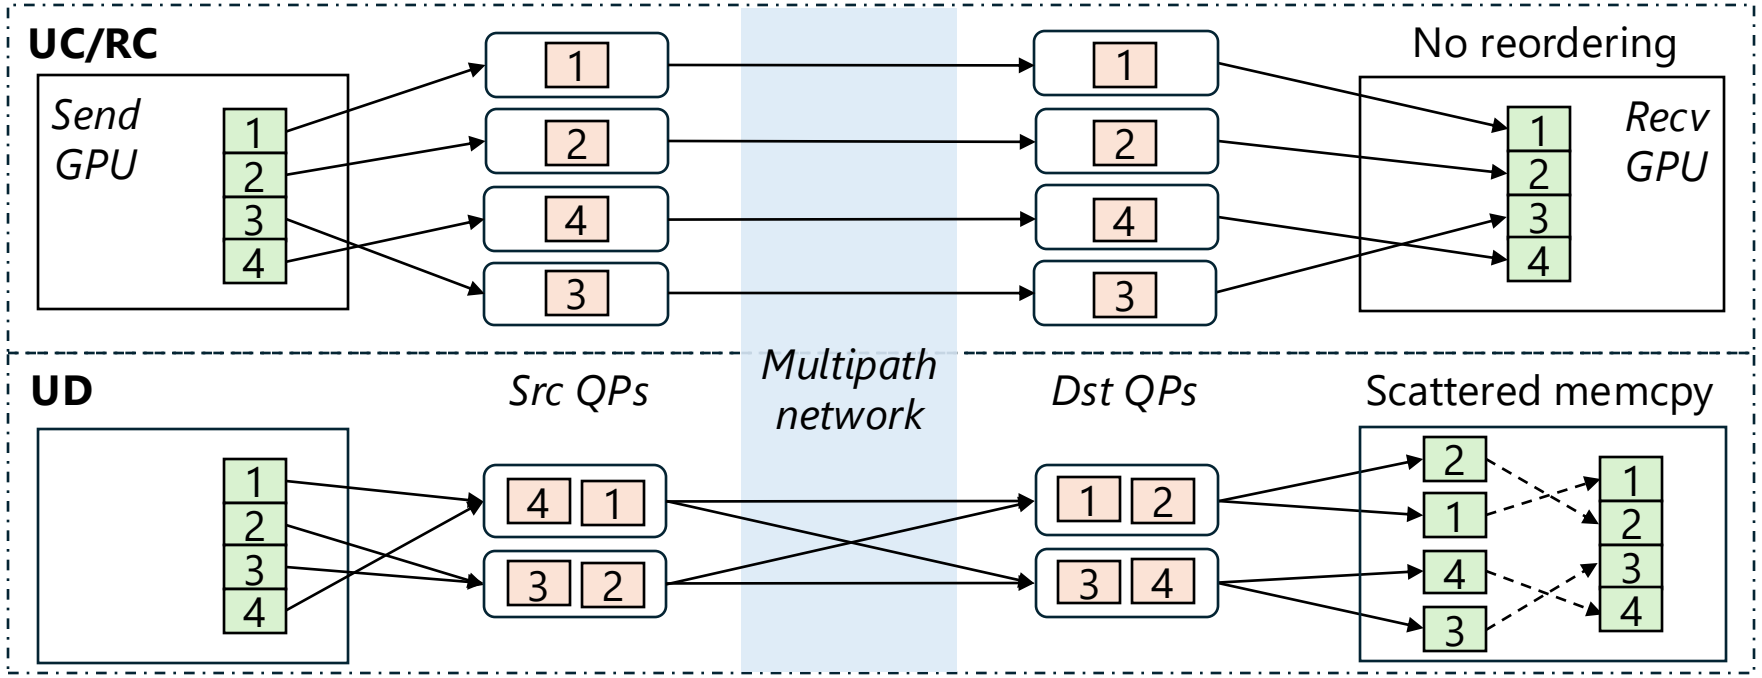

Technique #2: Efficient Software Transport

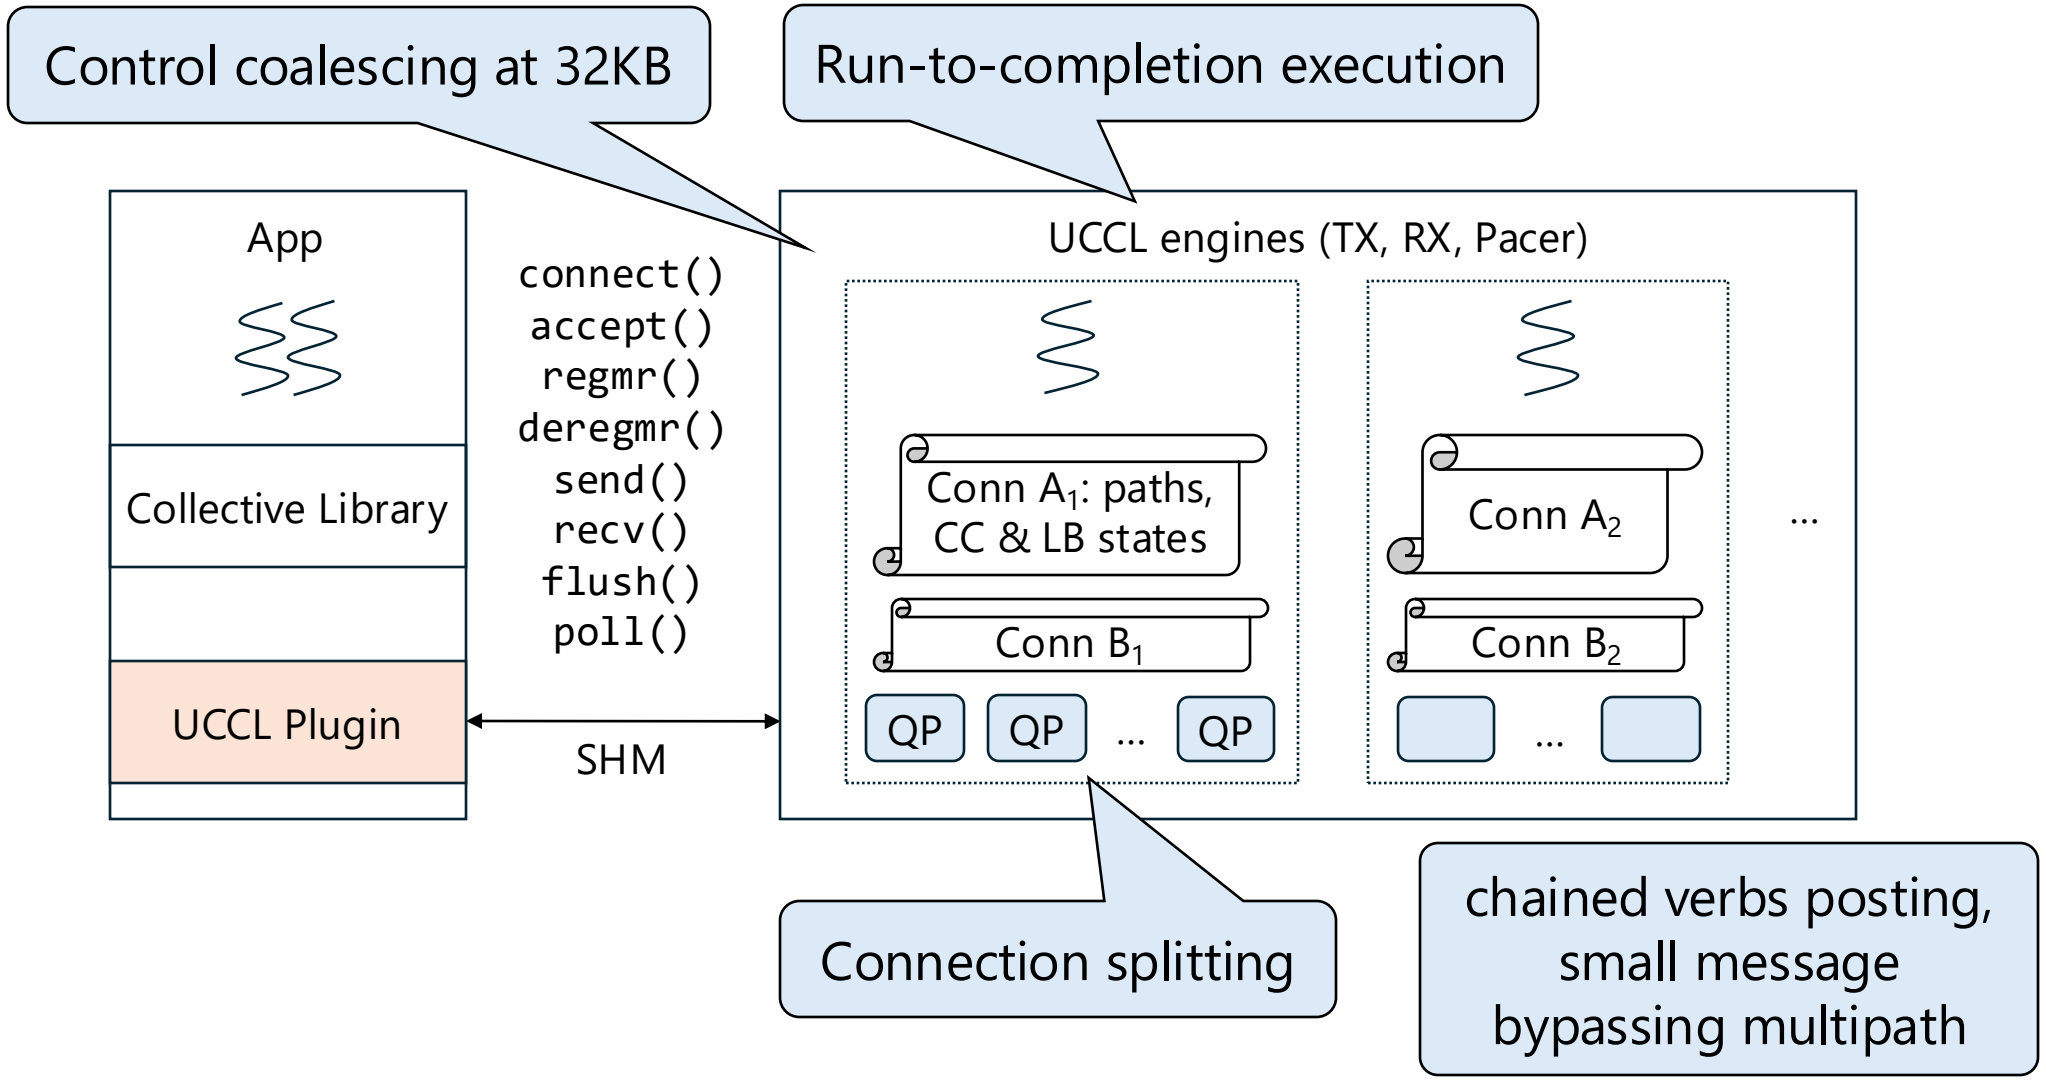

Implementation & Feature Support

- 27k LoC in C++
  - Drop-in replacement for NCCL applications
  - Packet spraying with 256 paths
  - Latency-based CC, receiver-driven CC
  - Efficient loss recovery by selective repeat
- Support both Nvidia and AMD GPUs
  - Future: AWS Trainium
- Support a variety of NIC vendors:
  - RDMA: Nvidia, Broadcom, AWS EFA
  - Non-RDMA: Nvidia, AWS ENA, IBM VirtIO

Evaluation: 4 AWS p4d all-to-all

- 4×100G EFA NICs per node, Fattree over Ethernet
  - NVLink disabled to emulate larger testbed

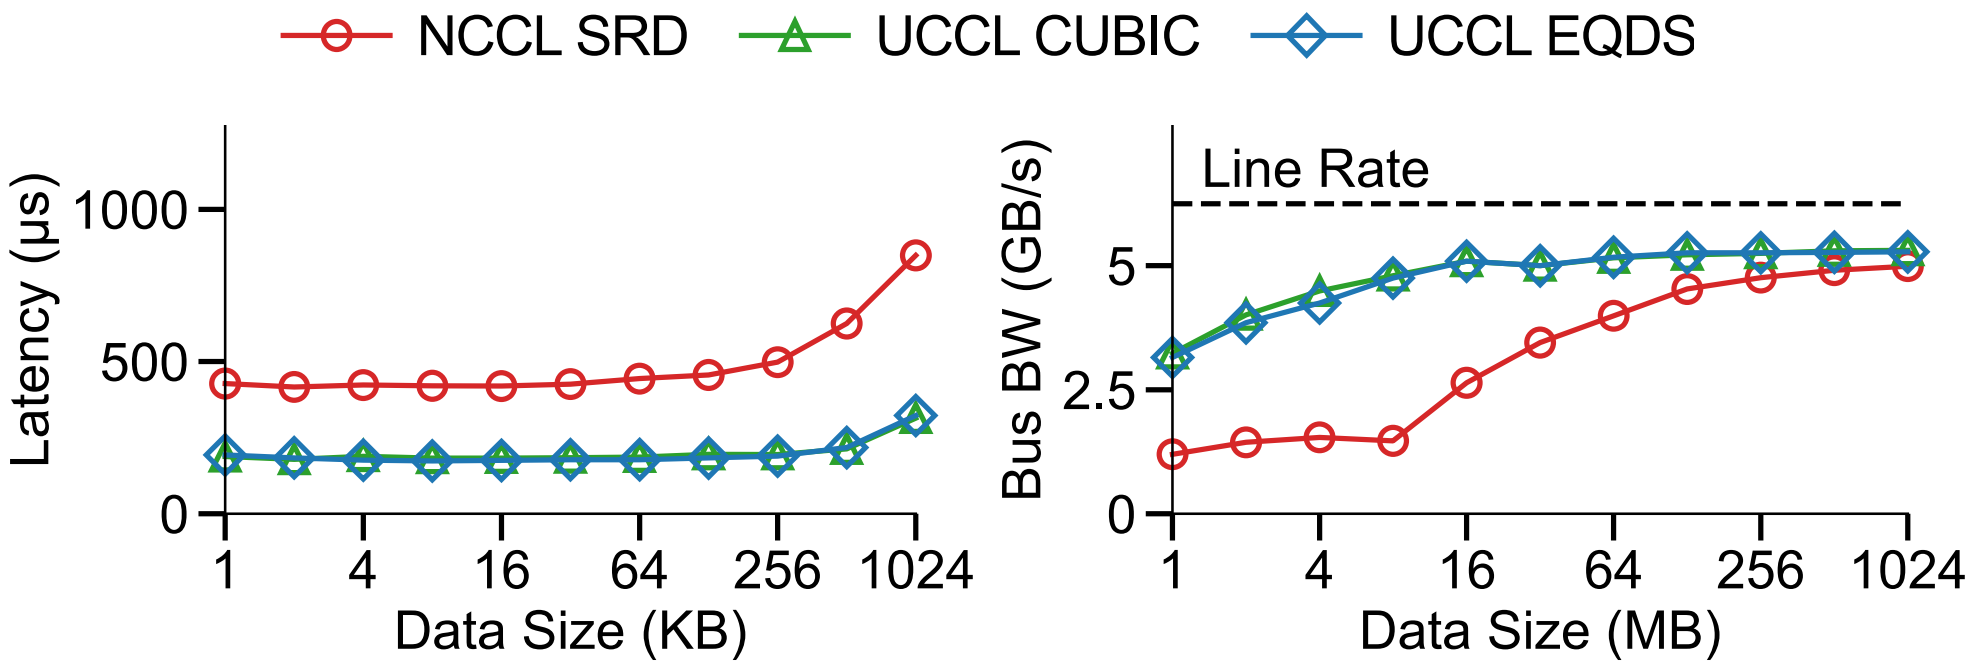

UCCL achieves up to 3.2× higher performance over NCCL on AWS

Evaluation: 2 HGX all-to-all same-rack IB

- 8×400G Nvidia CX-7 NICs per node, same rack over InfiniBand
  - NVLink disabled to emulate larger testbed

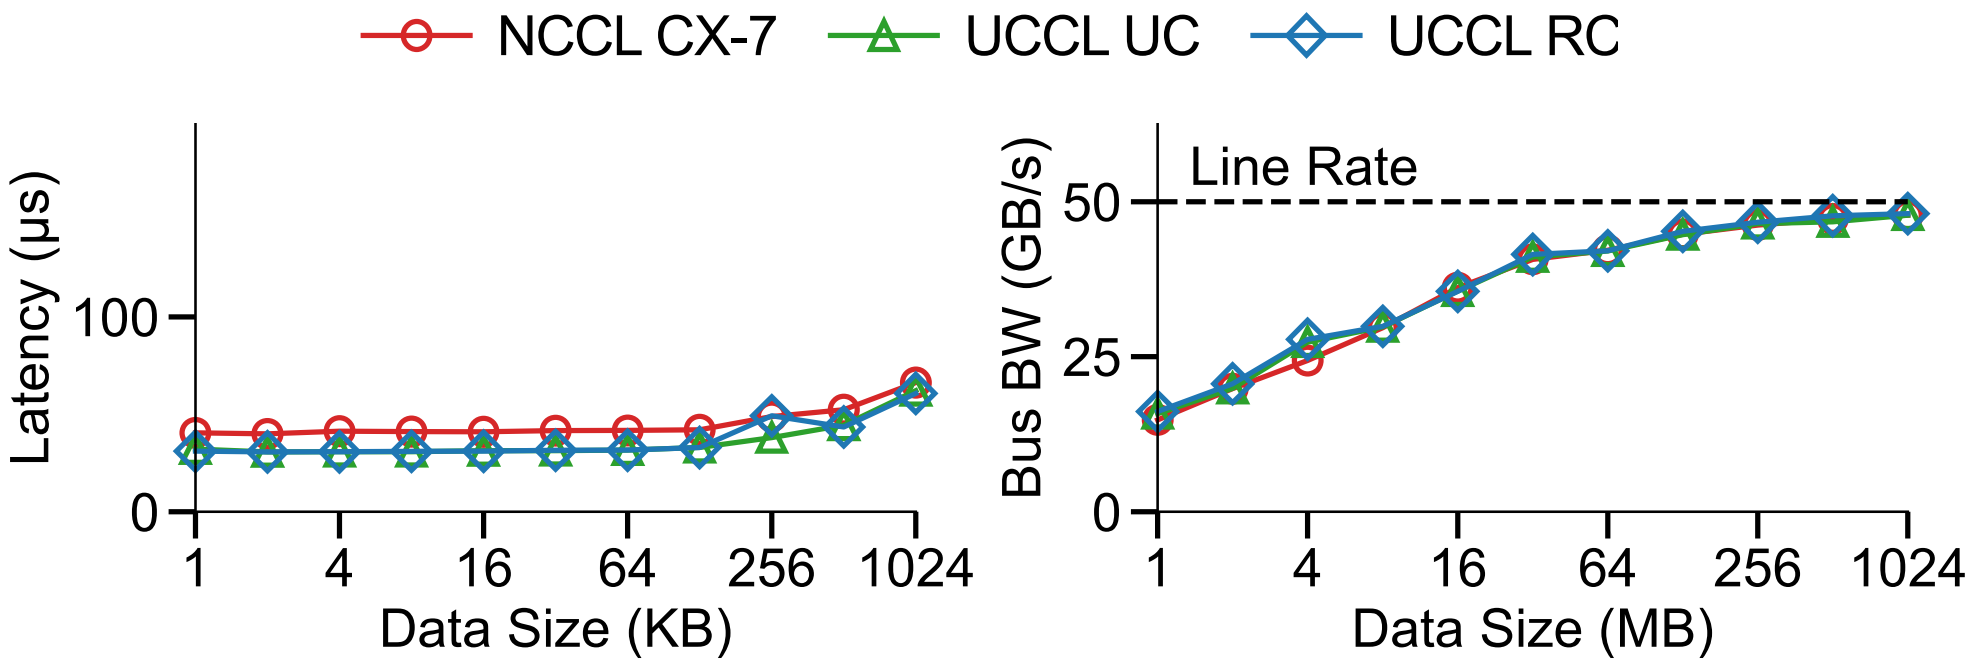

UCCL matches NCCL performances on ASIC-based NICs under same rack

Evaluation: 2 HGX all-to-all cross-rack RoCE

NCCL with RC QP scaling of 4

1024	16	float	none	-1	97.97	0.01	0.01
2048	32	float	none	-1	84.56	0.02	0.02
4096	64	float	none	-1	84.86	0.05	0.05
8192	128	float	none	-1	67.78	0.12	0.11
16384	256	float	none	-1	71.29	0.23	0.22
32768	512	float	none	-1	87.58	0.37	0.35
65536	1024	float	none	-1	82.16	0.80	0.75
131072	2048	float	none	-1	83.48	1.57	1.47
262144	4096	float	none	-1	75.96	3.45	3.24
524288	8192	float	none	-1	92.10	5.69	5.34
1048576	16384	float	none	-1	112.4	9.33	8.74
2097152	32768	float	none	-1	158.5	13.24	12.41
4194304	65536	float	none	-1	206.5	20.31	19.04
8388608	131072	float	none	-1	352.0	23.83	22.34
16777216	262144	float	none	-1	440.4	38.09	35.71
33554432	524288	float	none	-1	790.8	42.43	39.78
67108864	1048576	float	none	-1	1761.9	38.09	35.71
134217728	2097152	float	none	-1	3007.0	44.64	41.85
268435456	4194304	float	none	-1	5316.0	50.50	47.34
536870912	8388608	float	none	-1	10393	51.66	48.43
1073741824	16777216	float	none	-1	20644	52.01	48.76

Out of bounds values : 0 OK
Avg bus bandwidth : 17.8572
#

GB/s

UCCL with 256 paths (UC QPs)

1024	16	float	none	-1	92.63	0.01	0.01
2048	32	float	none	-1	93.10	0.02	0.02
4096	64	float	none	-1	82.55	0.05	0.05
8192	128	float	none	-1	85.56	0.10	0.09
16384	256	float	none	-1	89.27	0.18	0.17
32768	512	float	none	-1	78.68	0.42	0.39
65536	1024	float	none	-1	75.34	0.87	0.82
131072	2048	float	none	-1	72.96	1.80	1.68
262144	4096	float	none	-1	93.89	2.79	2.62
524288	8192	float	none	-1	111.8	4.69	4.40
1048576	16384	float	none	-1	105.0	9.99	9.36
2097152	32768	float	none	-1	136.2	15.40	14.44
4194304	65536	float	none	-1	185.8	22.57	21.16
8388608	131072	float	none	-1	285.3	29.40	27.56
16777216	262144	float	none	-1	360.7	46.52	43.61
33554432	524288	float	none	-1	533.5	62.89	58.96
67108864	1048576	float	none	-1	999.5	67.14	62.94
134217728	2097152	float	none	-1	1817.6	73.84	69.23
268435456	4194304	float	none	-1	3353.1	80.05	75.05
536870912	8388608	float	none	-1	6504.0	82.54	77.39
1073741824	16777216	float	none	-1	12851	83.5	78.33

Out of bounds values : 0 OK
Avg bus bandwidth : 26.156
#

GB/s

UCCL outperforms NCCL by 1.6× on ASIC-based NICs across racks

Dev Plan

- Dynamic membership with GPU servers joining and exiting
- GPU-initiated P2P communication (eg, IBGDA)
  - For MoE all-to-all and PD disaggregation
  - Generic to NIC vendors like AWS EFA and Broadcom, and GPU vendors like AMD
- Rearchitecting NCCL to unleash network hardware capability
  - Scalable and efficient CPU proxy
  - Low-cost async collectives with compute-communication ordering guarantee
  - Device kernels in vendor-agnostic Triton language

- We would like to hear about your feature needs!

Conclusion

UCCL: building the fastest collective communication library

- Network transport layer, CCL layer, heterogeneous vendors, and more
- Open and collaborative platform---talk with us in the poster session

tinyurl.com/uccl-paper

github.com/uccl-project/uccl

Thank you!

yangzhou.rpc@gmail.com/berkeley.edu